

# Graph Cut

## Second Milestone

Bruno Coswig Fiss    Max Göttner

<sup>1</sup>Institut für Technische Informatik und Mikroelektronik  
Technische Universität Berlin

Parallele Algorithmen auf GPUs



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Using Graph Cut for Energy Minimization

Minimize:

$$E(L) = \sum_x D_x(L_x) + \sum_{(x,y) \in N} V(|L_x - L_y|)$$

Sum over all Pixels of an Image

Sum over all neighborhoods

↓

↑

Data Term:  
Measures fitting of  
label to pixel

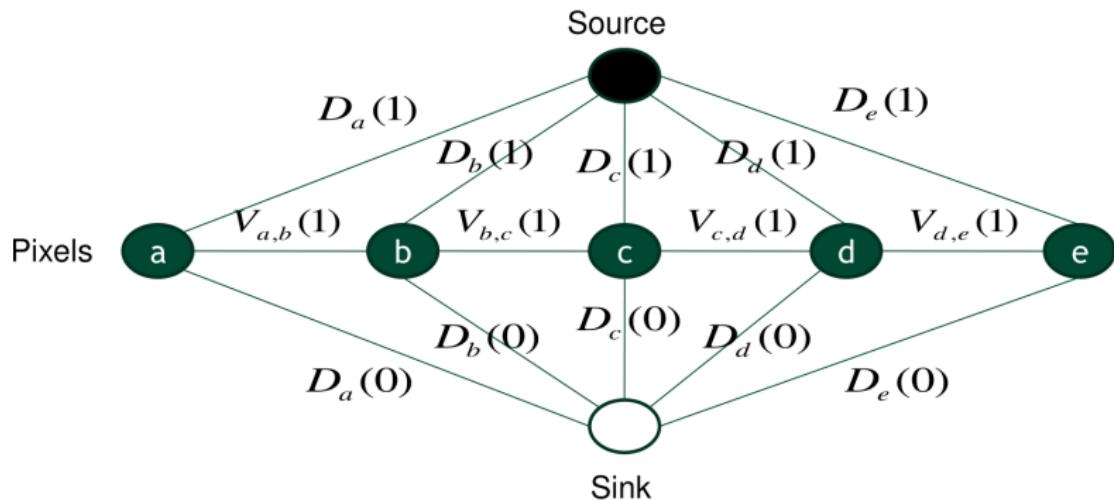
Neighborhood Term:  
Penalizes different labelings  
for neighbors

Credit: *Timo Stich, 2009*



# Using Graph Cut for Energy Minimization

Find minimum cut:

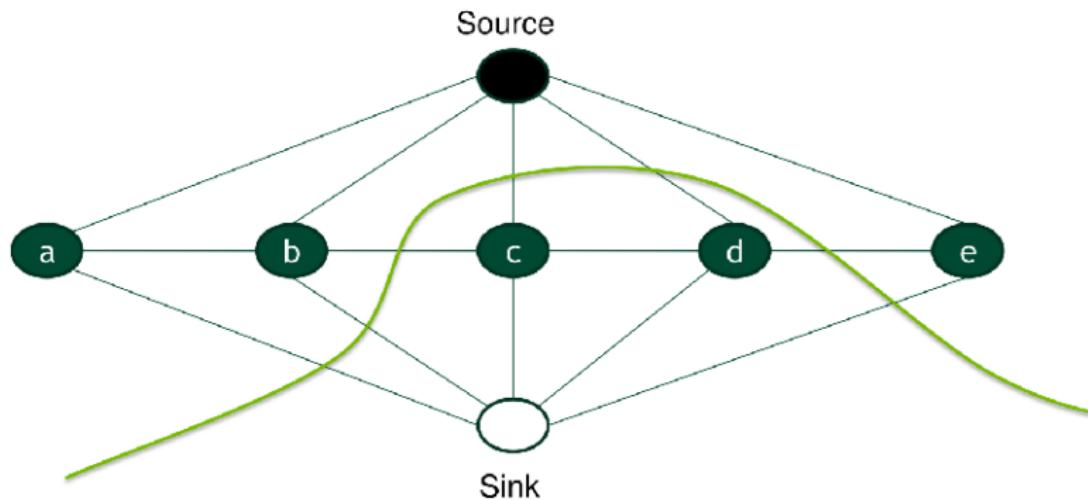


Credit: Timo Stich, 2009



# Using Graph Cut for Energy Minimization

Find minimum cut:



Credit: *Timo Stich, 2009*



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

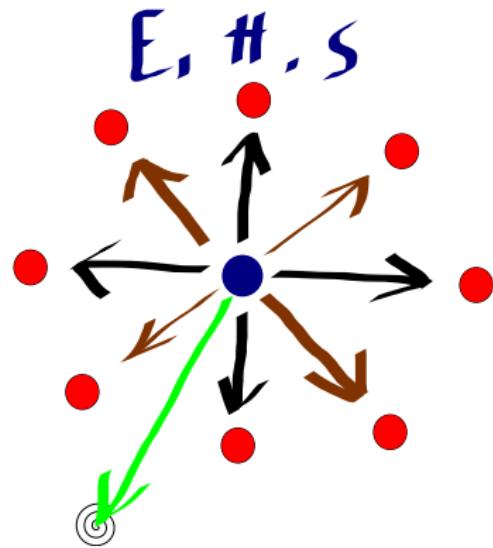
- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Information on Every Node



Every node has:

- Height.
- Excess flow.
- Status (speed-up).
- 4-8(!) edges, plus sink.

All are integers.



# Implementation Structure

- Simple idea: every thread is responsible for one node/pixel.
- Apply Push and Relabel to every node. Repeat it until flow stops changing.
- Size blocks appropriately (32x8 threads, based on previous works). Pad if necessary.
- One kernel to initialize graph.
- One kernel for Push and another for Relabel.
- Two kernels to extract labeling.



# Main Operations

Push:

- Read the height of all neighbors (**shared memory** applicable).
- Read/modify excess flow and edge capacities.  
Avoid hazards by using **atomic operations**.

Relabel:

- Read excess flow and edge information. **Boolean** would suffice.
- Read/modify heights (**shared memory** again applicable).



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

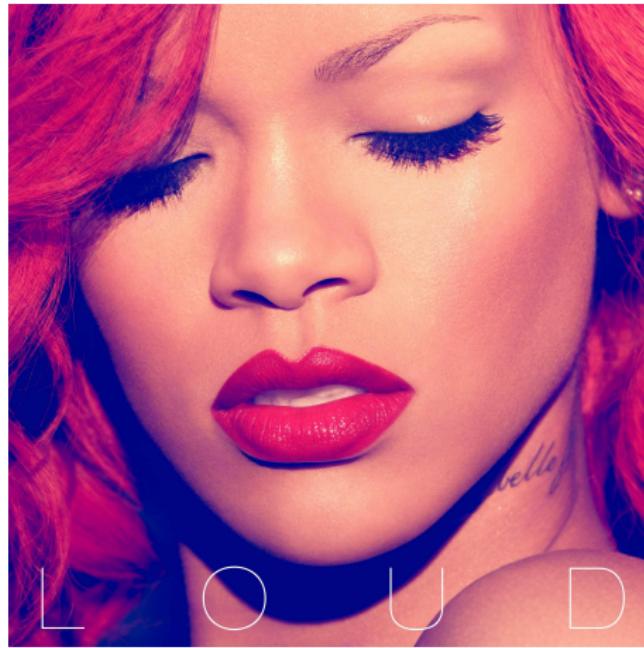
## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Neighborhood Comparison

Skin detection based on mixture of Gaussians by Jones, Regh.

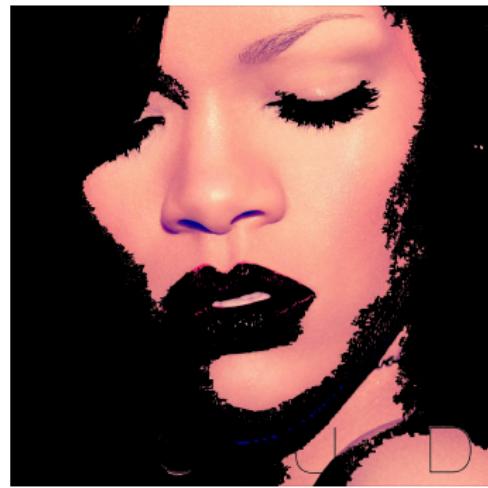


Original picture (3600x3600).

# Neighborhood Comparison



4-Neighborhood, 400 iterations.



8-Neighborhood, 4000 iterations.

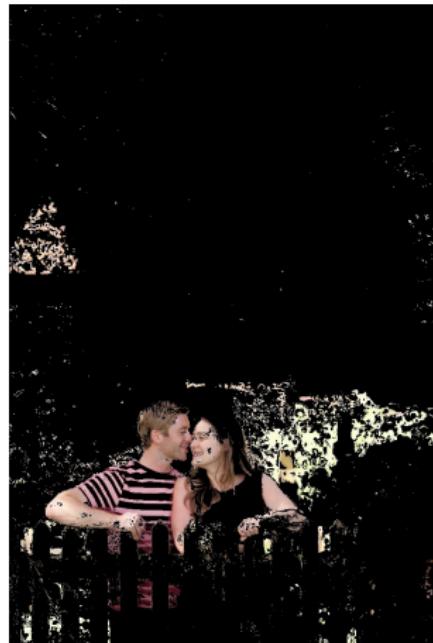


# More samples

Original and Naive



Original picture.

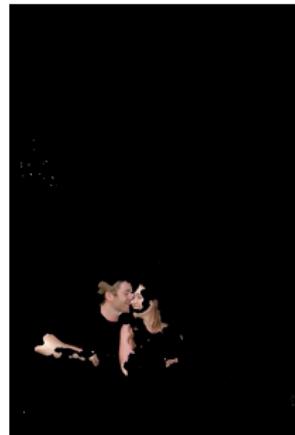
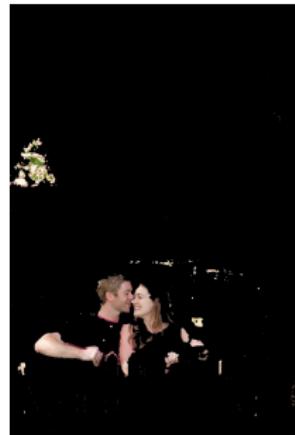
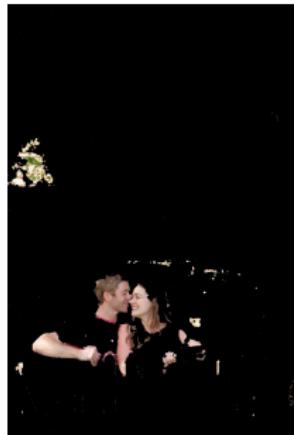


Naive segmentation.



# More samples

8-Neighborhood



Segmentation with  
 $P=1000$ .

Segmentation with  
 $P=1500$ .

Segmentation with  
 $P=2000$ .

Segmentation with  
 $P=2500$ .



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Comparison with Previous Works

Considered works include:

- CudaCuts by Vibhav Vineet and P. J. Narayanan, 2008.
- Own CPU implementation (for general graphs).
- CPU implementation (for specific application) by Boykov *et al.*.
- Timo Stich, 2009 (Nvidia ©, faster than CudaCuts).

Implementation	$\mu s$ per iteration
CudaCuts (640x480)	122.23
Ours with 8-N (640x480)	101.10
Ours with 4-N (640x480)	71.01
Ours with 4-N (3600x3600)	5319.82

# Further Comparison with Previous Works

- Currently only per iteration comparisons done.
- CPU implementations and GPU without source are not yet comparable.
- For a broader comparison, equal test cases must be used.
- Our number of iterations is still high(!): **Global Relabeling** technique necessary.



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# List of Performance Improvements

- Use shared memory for accessing heights (done).
- Keep list of active and inactive blocks (done, can be improved).
- Use Global Relabeling (currently “manually”).
- Height double buffering (because of memory boundness).
- Compress neighbor and edge status.
- Use Wave Push (and Wave Global Relabel).
- Use Dynamic Updating.

# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- **Problems**
- Other Applications



# List of Problems

- The implementation depends on atomic functions (sm\_11).
- Too many atomic accesses(?) end up generating: (6) the launch timed out and was terminated. [GPU also in use for display]
- Both can be solved by using the **Wave technique**.



# Outline

## 1 Our Implementation

- Recalling Problem...
- Algorithm Overview

## 2 Our Results

- Skin Detection
- Performance Comparison

## 3 Next Steps

- Performance Improvement
- Problems
- Other Applications



# Finding More Applications



Stereo Depth Estimation



Binary Image Segmentation



Photo Montage (aka Image Stitching)

Credit: MRF Evaluation,  
Middlebury College

- Skin detection (done - ~40ms per 800x600 frame).
- Segmentation background/foreground.
- Image Stitching.
- Stereo Depth Estimation.



# Conclusion

- Suggestions or questions?
- Thank you!
- Some additional material and references:
  - [http://cvit.iiit.ac.in/papers/rtGCuts\\_2008.pdf](http://cvit.iiit.ac.in/papers/rtGCuts_2008.pdf)
  - [http://www.nvidia.com/content/GTC/documents/1060\\_GTC09.pdf](http://www.nvidia.com/content/GTC/documents/1060_GTC09.pdf)
  - <http://vision.middlebury.edu/>